

DS 644: Homework 3

Instructions. Answer the following multiple choice questions by selecting the correct choices.

1. Principles of (functional) programming

- (a) Which of the following are programming paradigms? (Select three.)
- ☐ Declarative ☐ Functional ☐ Hadoop ☐ Imperative ☐ Scala
- (b) What three concepts characterize a purely functional programming language?
- ☐ referential transparency ☐ input/output ☐ no side effects ☐ procedural
 - ☐ immutability

2. Big data properties.

- (a) In lecture we discussed the meaning of the term “Big Data.” We decided that, for simplicity, we will call data “big” when it is
- ☐ at least 1Gb
 - ☐ too big to fit in fast memory (cpu cache + ram) on a single compute node
 - ☐ too big to fit in all computer memory (whether fast or slow)
 - ☐ too big to be dealt with by traditional data-processing software
- (b) “Big Data” concerns which of the following types of data?
- ☐ structured ☐ semi-structured ☐ unstructured ☐ all of these
- (c) JSON and XML are examples of which type of data?
- ☐ structured ☐ unstructured ☐ semi-structured ☐ none of these
- (d) Which two of the following statements are true of unstructured data?
- ☐ It is generally easier to analyze than other types of data.
 - ☐ It is often referred to as “messy” data.
 - ☐ It fits neatly into a schema.
 - ☐ It is the most widespread type of data.
 - ☐ It is usually found in tables.

3. Latency and fault-tolerance.

- (a) *Latency* is degradation in performance due to...
- ☐ a small number of cores in the central processing unit
 - ☐ slow data transfer across the network or cluster
 - ☐ shuffling data between different nodes in a cluster
 - ☐ failure of one or more nodes in the cluster
 - ☐ stack overflow caused by recursion
- (b) Hadoop achieves fault-tolerance by...
- ☐ using lazy evaluation and garbage collection.
 - ☐ writing intermediate computations to disk.
 - ☐ keeping all data immutable and in-memory.
 - ☐ replaying functional transformations over the original (immutable) dataset.

- (c) Spark decreases latency while remaining fault-tolerant by...
- ☐ using ideas from imperative programming.
 - ☐ using ideas from functional programming.
 - ☐ discarding data when it's no longer needed.
 - ☐ keeping all data immutable and in-memory.
 - ☐ replaying functional transformations over the original (immutable) dataset.

4. Transformations and actions.

- (a) In Spark a **transformation** on an RDD...
- ☐ is eagerly evaluated.
 - ☐ is lazily evaluated.
 - ☐ immediately computes and returns a result.
 - ☐ does not immediately compute a result.
 - ☐ usually returns another RDD (once it's evaluated).
- (b) In Spark an **action** on an RDD...
- ☐ is eagerly evaluated.
 - ☐ is lazily evaluated.
 - ☐ immediately computes and returns a result.
 - ☐ does not immediately compute a result.
 - ☐ always returns another RDD (once it's evaluated).
- (c) After performing a series of transformations on an RDD, which of the following methods could you use to make sure those transformations are not repeated (e.g., on each iteration of an algorithm)?
- ☐ `save`
 - ☐ `persist`
 - ☐ `memoize`
 - ☐ There is no such method because of the JVM's garbage collection mechanism.
- (d) Why does Spark's RDD class not have a `foldLeft` method?
- ☐ `foldLeft` can only be performed on lists of Boolean values.
 - ☐ `foldLeft` doesn't work on immutable collections.
 - ☐ `foldLeft` is not stack-safe.
 - ☐ `foldLeft` is not fault-tolerant.
 - ☐ `foldLeft` is not parallelizable.
- (e) Why is available in Spark's RDD class that overcomes the limitation of `foldLeft` mentioned in the previous part of this exercise?
- ☐ `aggregate` ☐ `fold` ☐ `foldLeft` ☐ `join` ☐ `leftOuterJoin`

5. **Read the docs.** Navigate to the Spark API documentation at

<https://spark.apache.org/docs/3.3.1/api/scala/org/apache/spark/index.html>

Enter “RDD” in the search box and select RDD from the results that appear on the left.

(a) Scroll down the resulting RDD API documentation page and find the `cache()` method. What does it say?

- ☐ Persist this RDD with the default storage level (`MEMORY_ONLY`).
- ☐ Mark the RDD as non-persistent, and remove all blocks for it from memory and disk.
- ☐ Set this RDD’s storage level to persist its values across operations after the first time it is computed.
- ☐ Save this RDD as a SequenceFile of serialized objects.

(b) On the RDD API doc page, find the version of `persist` that takes an argument:

`def persist(newLevel: StorageLevel)`. What does it say?

- ☐ Persist this RDD with the default storage level (`MEMORY_ONLY`).
- ☐ Mark the RDD as non-persistent, and remove all blocks for it from memory and disk.
- ☐ Set this RDD’s storage level to persist its values across operations after the first time it is computed.
- ☐ Save this RDD as a SequenceFile of serialized objects.

(c) On the RDD API doc page, find the `unpersist` method. What does it say?

- ☐ Persist this RDD with the default storage level (`MEMORY_ONLY`).
- ☐ Mark the RDD as non-persistent, and remove all blocks for it from memory and disk.
- ☐ Set this RDD’s storage level to persist its values across operations after the first time it is computed.
- ☐ Save this RDD as a SequenceFile of serialized objects.

(d) What’s the difference between the `sample` and `takeSample` methods of the RDD class?

- ☐ `sample` always uses a with-replacement sampling method, while `takeSample` always samples without replacement.
- ☐ `sample` returns an RDD, while `takeSample` returns an Array.
- ☐ The second argument specifies the number of samples desired either as a fraction of the size of the RDD (`sample`) or as an absolute number (`takeSample`).
- ☐ There is no difference; they are just two different names one can use to invoke the same function.